

Sardines: A Networked Game

Sardines is a system of communication for 2-8 players. Each player pilots a submarine in real time, viewing their surroundings through the lens of a retro sonar system. These submarines communicate in morse code, firing soundwaves - representative of dots and dashes - to one another across the map. While this system has ultimately been designed for integration into a larger, more complete game, *Sardines* presents itself as a sandbox, to best restrict attention to the networking techniques at play.

Architecture

Sardines' network uses a straightforward client-server architecture with local input processing - that is to say, each client makes [

Bauer et al. (2004) evaluate the scalability of this architecture, finding that with n entity states (for the purposes of *Sardines*, players), client-server costs grow at order $O(n^2)$, compared to peer-to-peer's $O(n^3)$. They ultimately conclude that "The client-server architecture exhibits the lowest growth in overall system cost, however, with the disadvantage that the entire growth must be handled by the central server."

In light of their findings, this report proposes a hybrid architecture for *Sardines* to adopt at scale. While it ultimately proved too ambitious for this assignment, the original idea for . Since only navigators would be able to see the surrounding world, they could in some ways act as middle men in the hybrid architecture set out in Figure [FIGURE]. By treating each navigator as a 'local' server for their crew of 4, ... This would, of course, be subject to further testing, as [carries issues of P2P?]... [In paragraph/figure, introduce idea of navigator 'standardising' actions of the crew...]

[GRAPHICSX: Handdrawn diagram of interpolation and prediction interaction?]

Protocols

Transport Layer TCP: Reliable, etc... As much is set out in RFC 798 , where John Postel introduces the protocol:.... Provide reliability paragraph?...

There is, of course, an argument for updating positions via UDP. Suppose that a submarine starts at position *A*, fails to update to position *B*: where the ordered nature of TCP means the application will resend and resend this update [EXPLAIN DOWNSIDE], this protocol will simply [GO TO the position C of the next successful]. While the deliberate simplicity of UDP - the [FEATURE #1], the [FEATURE #2] - lends itself to the continuous, incremental nature of movement , *Sardines* will not make use of it. Submarines in this game travel slowly, with positions needed reliably but not immediately, and with the robustness of the prediction system used (see 'Prediction'), updating position via TCP every 0.1s should be infrequent enough to avoid the backlog described above.¹

Application Layer While it doesn't consider the internetwork and hardware layers of the protocol stack, the application layer of *Sardines*' network necessarily interacts with the transport layer directly below it. Client and server communicate using an underlying **Packet** struct, ... [Serialisation!]...

At a more granular level, *Sardines*' **SendablePackets** may use any such struct as a **serialisedBody**:

- **SyncPacket** Contains a long **syncTimestamp**. Sent on connection to standardise server and client calculations for **DateTime.UtcNow.Ticks** (which may vary from underlying OS to underlying OS, Microsoft n.d.a).
- **IDPacket** ... (in the hybrid architecture discussed above, this IP would also be used to establish crew-to-navigator connections).
- **SubmarinePacket**
- **PositionPacket**
- **MorsePacket**
- **EmptyPacket** Contains no variables. Sent when the **bodyID** corresponds to a function with no arguments (e.g. when a client starts a game with **bodyID**).

¹Could simplify even further with event-based - have server calculate all positions from key presses...

bodyIDs use the following naming convention: IDs 1XXX correspond to clients connecting to/disconnecting from a server, 2XXX to server functionality while in lobby mode, 3XXX to server functionality while in a game mode, and 4XXX to client actions which in-game.

While there isn't the space to break down every protocol in precise detail, consider the process of a player joining a lobby:

1. *The client registers a TCP connection with the server.* The client then constructs... with ID 1000...
2. *The server receives a **SyncPacket** from the client.* Calling `Receive1000()`, the server...
3. *The client receives a **SyncPacket** from the server.*
4. *The server receives an **IDPacket** from the client.*
5. *The client receives an **IDPacket** from the server.*
6. *All clients receive (further) **IDPackets** from the server.*

There is arguably some unnecessary back-and-forth to the above, but small packet sizes shouldn't put any meaningful strain on bandwidth. The process is designed first and foremost for ease of programming: treating major protocols as a chain of smaller, simpler steps, it becomes far easier to manage - and document - the application layer.

API

Sardines is built with C# in the Godot engine. It uses `System.Net.Sockets` to handle networking, and `System.Runtime.InteropServices` to serialize/deserialize packet structs. As noted in Microsoft's documentation (n.d.b), `System.Net.Sockets` implements conventional Berkeley sockets.

Integration

Asynchronous I/O... Connection class...

Discuss: offline vs. online updates to position!

This report would be amiss to skip over the final step:

Prediction

As discussed under 'Architecture', clients only send position updates every 0.1s. What this report has so far failed to consider is how this appears to other clients - they experience what should be a smooth, continuous movement as discontinuous jumps over 0.1s intervals! Clearly, ... [introduce prediction - with reading?].

When a player chooses to move forward, they do not jump to a constant speed but continuously accelerate from zero; naturally, *Sardines* uses second-order quadratic prediction to best approximate the second-order derivative of acceleration. Given a submarine's three most recent positions $\mathbf{r}_0, \mathbf{r}_1, \mathbf{r}_2$ (corresponding to times $t_0 > t_1 > t_2$), clients can average the velocities from \mathbf{r}_1 to \mathbf{r}_0 , from \mathbf{r}_2 to \mathbf{r}_1 , and the acceleration from \mathbf{r}_1 to \mathbf{r}_0 as

$$\mathbf{u}_0 = \frac{\mathbf{r}_0 - \mathbf{r}_1}{t_0 - t_1}, \quad \mathbf{u}_1 = \frac{\mathbf{r}_1 - \mathbf{r}_2}{t_1 - t_2}, \quad \mathbf{a}_0 = \frac{\mathbf{u}_0 - \mathbf{u}_1}{t_0 - t_1}, \quad \text{respectively.}$$

These estimates define the quadratic model

$$\tilde{\mathbf{r}}(t) = \mathbf{r}_0 + \mathbf{u}_0 t + \mathbf{a}_0 t^2.$$

In contrast, the rudder controlling a submarine's rotation θ is controlled at a constant speed, so *Sardines* only uses linear prediction to approximate

$$\tilde{\theta}(t) = \theta_0 + \dot{\theta}_0 t$$

(with subtle, case-specific considerations made given $\theta \in [0, 2\pi)$).

If prediction is the act of waiting for data, then integration is how one 'catches up' on receiving it. On receiving a new `PositionPacket` at time t_0 , a programmer might be inclined to start predicting under to a new quadratic model $\tilde{\mathbf{r}}_{\text{new}}(t)$ immediately, but if positions $\tilde{\mathbf{r}}_{\text{old}}(t_0)$ and $\tilde{\mathbf{r}}_{\text{new}}(t_0)$ are visibly far apart, then the

player will see the corresponding submarine make an instantaneous jump across the screen.² Instead, one takes a set time T to linearly interpolate from the old trajectory to the new:

$$\tilde{\mathbf{r}}(t) = \begin{cases} \tilde{\mathbf{r}}_{\text{old}}(t) & \text{if } t < t_0 \\ (1 - q(t))\tilde{\mathbf{r}}_{\text{old}}(t) + q(t)\tilde{\mathbf{r}}_{\text{new}}(t) & \text{if } t_0 \leq t < t_0 + T, \text{ where } q(t) = \frac{1}{T}(t - t_0). \\ \tilde{\mathbf{r}}_{\text{new}}(t) & \text{if } t \geq t_0 + T \end{cases}$$

In *Sardines*' particular implementation, **PositionPackets** are sent via TCP every 0.1s; interpolation therefore takes place over a strictly shorter interval $T = 0.05s$.

[GRAPHICSX: Handdrawn diagram of interpolation and prediction interaction?]

To fully understand how *Sardines* uses its prediction techniques, this report must first introduce a core challenge of any networked game: conflict resolution.

In *Sardines*, the projectiles concerned are soundwaves. The visual language of the game, where soundwaves from external sources only become visible on collision with the player, provides a clear approach: the sender unequivocally takes precedence. Only when a player sees their soundwave hit another is a **MorsePacket** sent from their client (which will arrive with the usual delay). The sender knows with certainty who receives their message; the receiver, who cannot see the trajectory of the soundwave until it arrives, will have no sense of whether it "should" have hit them.

To further 'smooth over' the application's conflict resolution, the receiving client makes use of backward prediction. Since neither server nor client stores more than three of any submarine's past positions at a time, it is fortunate the above formulae can approximate the past as well as the future.³

Suppose a sender emits a soundwave from position \mathbf{r} at time t_0 , which they see reach a receiver at $t_0 + \Delta t$. On the arrival of the corresponding packet at t_1 , then, the receiving client has to decide where the wave was emitted from *in its local view of the game*. The obvious choice would be the 'true origin' \mathbf{r} , but *Sardines* uses the backwards prediction $\tilde{\mathbf{r}}(t_1 - \Delta t)$. As [FIGURE] puts it in [REFERENCE], [QUOTE]; conflict resolution is the art of deciding which quantities are preserved across clients, and *Sardines* - a system designed around slow, real-time communications - is far less concerned with a shared view of geography than it is a shared view of delay.

Testing

[SORT THIS LAST THING - BUT PLAN THE TESTING OUT BY 15th?]

References

- Bauer, D., Iliadis, I., Rooney, S. & Scotton, P. (2004), 'Communication architectures for massive multi-player games', *Multimedia Tools and Applications* **23**(1), 47–66.
- Microsoft (n.d.a), 'Learn / .NET / .NET API browser / System / DateTime / Properties / DateTime.UtcNow Property', Available at: https://learn.microsoft.com/en-us/dotnet/api/system.datetime.utcnow?redirectedfrom=MSDN&view=net-7.0#System_DateTime_UtcNow. (Accessed: 14 January 2023).
- Microsoft (n.d.b), 'Learn / .NET / .NET API browser / System.Net.Sockets Namespace', Available at: <https://learn.microsoft.com/en-us/dotnet/api/system.net.sockets?view=net-7.0>. (Accessed: 14 January 2023).

²This might be regarded interpolation over $T = 0.0s$!

³*Sardines*' submarines are physics-based objects, and at one point in development, the drag they experience was factored into prediction. However, the differential equations for 2D motion with a quadratic drag were too complex to find an analytic solution - rather than being able to substitute a t -value into a given equation, the prediction would be calculated over incremental, irreversible forward time steps - so the application sacrifices this more realistic model for the ability to look backwards in time.